




🔑 NEEDS A CLOUD API KEY

- LLM — recommended cloud (OpenAI / Gemini) for demo latency (local HF model is the zero-key fallback)
- Polly TTS — only on the Polly viseme path (see viseme diagram)
- Deepgram / Google ASR — optional swap for Whisper (local, no key)

Everything else runs local / CPU / no key.